



The EigenTrust Algorithm for Reputation Management in P2P Networks

Sepandar D. Kamvar

Mario T. Schlosser

Hector Garcia-Molina

Stanford University


P2P Networks

- Open and anonymous
 - Benefits
 - Robust, Scalable, Diverse
 - Problems
 - Malicious peers
 - Inauthentic files
 - Viruses/Malware
 - Tampered files

Identifying malicious peers is more pressing than identifying inauthentic files

Reputation Systems


- Reputation Systems
 - *Global*: Centralized system (eBay)
 - *Local*: Distributed System
- *Key Idea of EigenTrust*: Each peer i is assigned a *global* trust value that reflects the *local* experiences of all the peers in the network



Problem

- **Problem:**

- Reduce inauthentic files distributed by malicious peers on a P2P network.




- **Motivation:**

“Major record labels have launched an aggressive new guerrilla assault on the underground music networks, flooding online swapping services with bogus copies of popular songs.”

-Silicon Valley Weekly

Problem

- **Goal:** To identify sources of inauthentic files and bias peers against downloading from them.
- **Method:** Give each peer a *trust value* based on its previous behavior.




Some approaches

- Past History
- Friends of Friends
- EigenTrust


Terminology

- **Local trust value: c_{ij} .**
The opinion that peer i has of peer j , based on past experience.
- **Global trust value: t_i .**
The trust that the entire system places in peer i .





Local Trust Values

- Each time peer i downloads an authentic file from peer j , c_{ij} increases.
- Each time peer i downloads an inauthentic file from peer j , c_{ij} decreases.




Normalizing Local Trust Values

- All c_{ij} non-negative
- $c_{i1} + c_{i2} + \dots + c_{in} = 1$



Local Trust Vector

- **Local trust vector \mathbf{c}_i :**
contains all local trust values c_{ij} that peer i has of other peers j .



$$\begin{pmatrix} 0 \\ \text{Peer 2} \\ 0 \\ \text{Peer 4} \\ \text{Peer 1} \end{pmatrix} = \begin{pmatrix} 0 \\ 0.9 \\ 0 \\ 0.1 \end{pmatrix} \mathbf{c}_1$$

Approach 1: Past history


- Each peer biases its choice of downloads using its own opinion vector \mathbf{c}_i .
- If it has had good past experience with peer j , it will be more likely to download from that peer.
- **Problem:** Each peer has limited past experience. Knows few other peers.



Approach 2: Friends of Friends

- Ask for the opinions of the people who you trust.

(Cf. Referral trust)




Friends of Friends

- Weight their opinions by your trust in them.

(Cf. Referral trust = Functional trust)

(Cf. Transitivity)



The Math : Transitive Trust

$$c'_{ik} = \sum_j c_{ij} \cdot c_{jk}$$

← What they think of peer k.

Ask your friends j

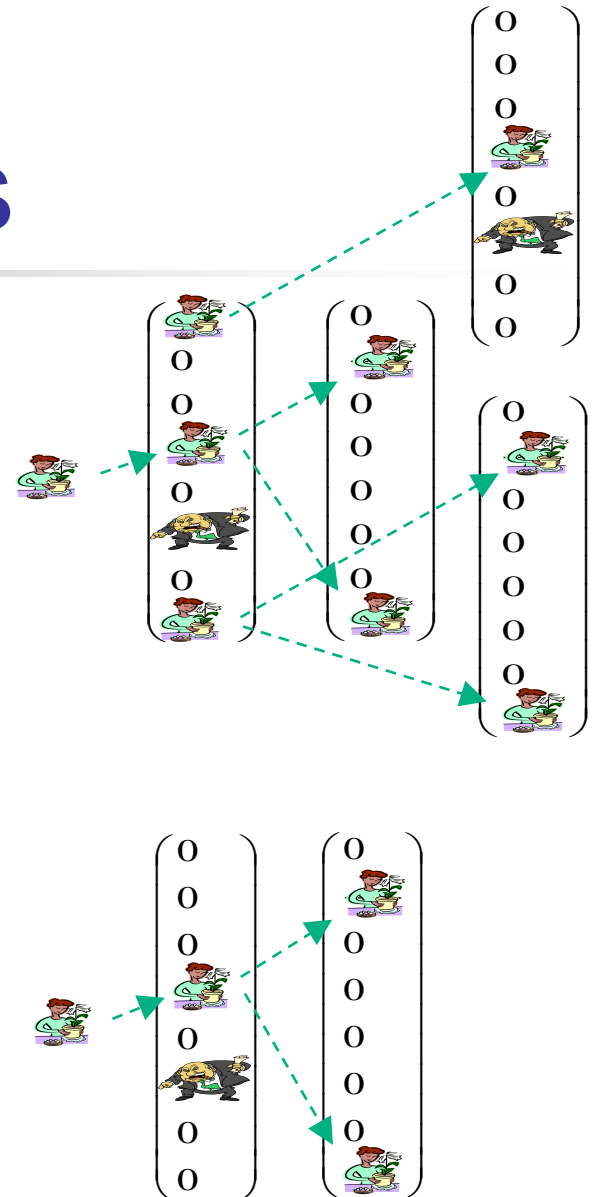
And weight each friend's opinion by how much you trust him.

$$\mathbf{c}'_i = \mathbf{C}^T \mathbf{c}_i$$

The diagram illustrates the matrix multiplication $\mathbf{c}'_i = \mathbf{C}^T \mathbf{c}_i$. On the left, a blue column vector \mathbf{c}'_i contains the values [.1, .3, .2, .3, .1, .1]. In the center, a green matrix \mathbf{C}^T is shown, with its second row highlighted in red and containing the values [0.2, 0.3, 0, 0.5, 0.1, 0, 0, 0]. On the right, a blue column vector \mathbf{c}_i contains the values [.1, .5, 0, 0, 0, .2]. The equation is represented by an equals sign between the vector and the matrix product.

Problem with Friends

- Either you know a lot of friends, in which case, you have to communicate, compute and store many values.
- Or, you have few friends, in which case you won't know many peers, even after asking your friends.




Eigen Trust: Dual Goal

- We want each peer to:
 - Know all peers.
 - Perform minimal computation (and storage).

Knowing All Peers

- Ask your friends:
 $t = C^T c_i$.
- Ask their friends:
 $t = (C^T)^2 c_i$.
- Keep asking until
the cows come
home: $t = (C^T)^n c_i$.



Minimal Computation

- Luckily, the *trust vector* \mathbf{t} , if computed in this manner, converges to the same thing for every peer!
- Therefore, each peer doesn't have to store and compute its own trust vector. The whole network can cooperate to store and compute \mathbf{t} .

Non-distributed Algorithm

- Initialize:

$$\mathbf{t}^{(0)} = \begin{bmatrix} \frac{1}{n} & \dots & \frac{1}{n} \end{bmatrix}^T$$

- Repeat until convergence:

$$\mathbf{t}^{(k+1)} = \mathbf{C}^T \mathbf{t}^{(k)}$$

Basic EigenTrust Algorithm

- Assumption: include central server at this stage
 - A server stores all the c_{ij} values and performs the computation

$$\vec{t}^{(0)} = \vec{e}; \quad e_i = 1/n$$


repeat

$$\vec{t}^{(k+1)} = C^T \vec{t}^{(k)};$$


$$\delta = \|\vec{t}^{(k+1)} - \vec{t}^{(k)}\|;$$

until $\delta < \epsilon$;

An Illustration Example of EigenTrust



Normalization

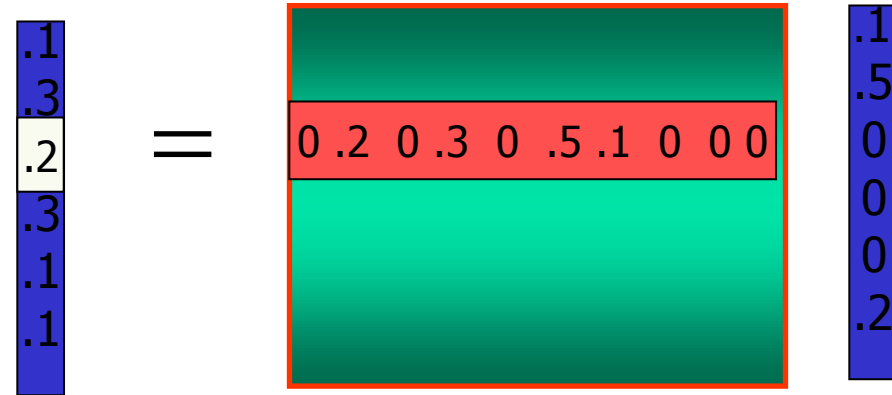


$$C = \begin{bmatrix} 0 & 0.67 & 0.33 & 0 & 0 & 0 \\ 0.5 & 0 & 0 & 0 & 0 & 0.5 \\ 0.08 & 0 & 0 & 0.23 & 0.69 & 0 \\ 0 & 0.78 & 0 & 0 & 0 & 0.22 \\ 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0.67 & 0 & 0 & 0.33 & 0 \end{bmatrix}$$

$$\begin{array}{ccccc} \begin{matrix} 0.1667 \\ 0.1667 \\ 0.1667 \\ 0.1667 \\ 0.1667 \end{matrix} & \begin{matrix} 0.0967 \\ 0.3534 \\ 0.0550 \\ 0.2050 \\ 0.1700 \\ 0.1200 \end{matrix} & \begin{matrix} 0.1811 \\ 0.3051 \\ 0.0319 \\ 0.1827 \\ 0.0776 \\ 0.2218 \end{matrix} & \begin{matrix} 0.1764 \\ 0.3434 \\ 0.0582 \\ 0.1188 \\ 0.1055 \\ 0.1979 \end{matrix} \\ t^0 = & t^1 = C^T t^0 = & t^2 = C^T t^1 = & \dots\dots\dots \end{array}$$

Distributed Algorithm

- No central authority to store and compute \mathbf{t} .
- Each peer i holds its own opinions \mathbf{c}_i .
- For now, let's ignore questions of lying, and let each peer store and compute its own trust value.



$$t_i^{(k+1)} = c_{1i} t_1^{(k)} + \dots + c_{ni} t_n^{(k)}$$

Distributed Algorithm

For each peer i {

- First, ask peers who know you for their opinions of you.

- Repeat until convergence {

 - Compute** current trust value: $t_i^{(k+1)} = c_{1i} t_1^{(k)} + \dots + c_{ni} t_n^{(k)}$


 - Send** your opinion c_{ij} and trust value $t_i^{(k)}$ to your acquaintances.

 - Wait** for the peers who know you to send you their trust values and opinions.

 - }


}

Probabilistic Interpretation : Random Surfer Model



Malicious Collectives :

Random Jumps to avoid dead ends



Revised Non-distributed Algorithm

- Initialize:


$$\mathbf{t}^{(0)} = \begin{bmatrix} \frac{1}{n} & \dots & \frac{1}{n} \end{bmatrix}^T$$

- Repeat until convergence:

$$\mathbf{t}^{(k+1)} = a\mathbf{C}^T\mathbf{t}^{(k)} + (1-a)\mathbf{p}$$

Pre-trusted Peers

- Battling Malicious Collectives
- Inactive Peers
- Incorporating heuristic notions of trust
- Improving Convergence Rate



Practical Issues

- Apriori notions of trust
 - Can we assign any profit to newcomers?
 - Only the first few peers to join the network are known to be trustworthy
 - $p_i = 1/|P|$ if $i \in P$, and $p_i = 0$ otherwise
 - Use \bar{p} instead of \bar{e}

Practical Issues(2)

- Inactive Peers
 - What happens if peer i doesn't download from anybody else?

$$c_{ij} = \begin{cases} \frac{\max(s_{ij}, 0)}{\sum_j \max(s_{ij}, 0)} & \text{if } \sum_j \max(s_{ij}, 0) \neq 0; \\ p_j & \text{otherwise} \end{cases}$$

- Choose to trust the pre-trusted peers

Practical Issues(3)

- Malicious Collectives

- a group of malicious peers who know each other
- How to prevent them from subverting the system?

$$\mathbf{t}^{(k+1)} = (1-a)C^T \mathbf{t}^{(k)} + a\mathbf{p}$$

- The modified algorithm:

$$\vec{t}^{(0)} = \vec{p};$$

repeat

$$\left| \begin{array}{l} \vec{t}^{(k+1)} = C^T \vec{t}^{(k)}; \\ \vec{t}^{(k+1)} = (1-a)\vec{t}^{(k+1)} + a\vec{p}; \\ \delta = ||\vec{t}^{(k+1)} - \vec{t}^{(k)}||; \end{array} \right.$$

until $\delta < \epsilon$;

Distributed EigenTrust

- Assumption: Everyone is honest
- Each peer computes its own global trust value:

$$t_i^{(k+1)} = (1-a)(c_{1i}t_1^{(k)} + \dots + c_{ni}t_n^{(k)}) + ap_i$$

Definitions:

- A_i : set of peers which have downloaded files from peer i
- B_i : set of peers from which peer i has downloaded files

Algorithm:

Each peer i do {

Query all peers $j \in A_i$ for $t_j^{(0)} = p_j$;

repeat

 Compute $t_i^{(k+1)} = (1-a)(c_{1i}t_1^{(k)} + c_{2i}t_2^{(k)} + \dots + c_{ni}t_n^{(k)}) + ap_i$;

 Send $c_{ij}t_i^{(k+1)}$ to all peers $j \in B_i$;

 Compute $\delta = |t_i^{(k+1)} - t_i^{(k)}|$;


 Wait for all peers $j \in A_i$ to return $c_{ji}t_j^{(k+1)}$;

until $\delta < \epsilon$;

}


Algorithm Complexity

- The algorithm converges fast
 - A network of 100 peers after 100 query cycles




Secure Score Management

- Two basic ideas:
 - Instead of having a peer compute and store its own score, *have another peer compute and store its score.*
 - *Have multiple score managers who vote on a peer's score.*



Distributed Hash Table




How to use the trust values t_i

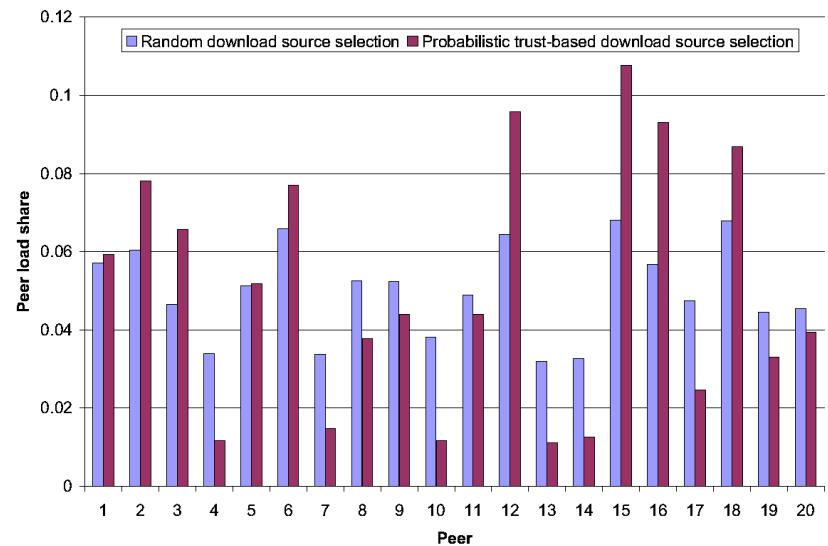
- When you get responses from multiple peers:
 - **Deterministic:** Choose the peer with highest trust value.
 - To avoid discriminating against new peer (at the risk of inviting malicious peer), occasionally use other peer.
 - **Probabilistic:** Choose a peer with probability proportional to its trust value.
 - This approach improves *load balancing*.

Load Distribution

Deterministic Download Choice




Probabilistic Download Choice



Threat Scenarios


■ Malicious Individuals

- Always provide inauthentic files.



■ Malicious Collective


- Always provide inauthentic files.
- Know each other. Give each other good opinions, and give other peers bad opinions.



More Threat Scenarios

■ Camouflaged Collective

- Provide authentic files some of the time to trick good peers into giving them good opinions.




■ Malicious Spies


- Some members of the collective give good files all the time, but give good opinions to malicious peers.



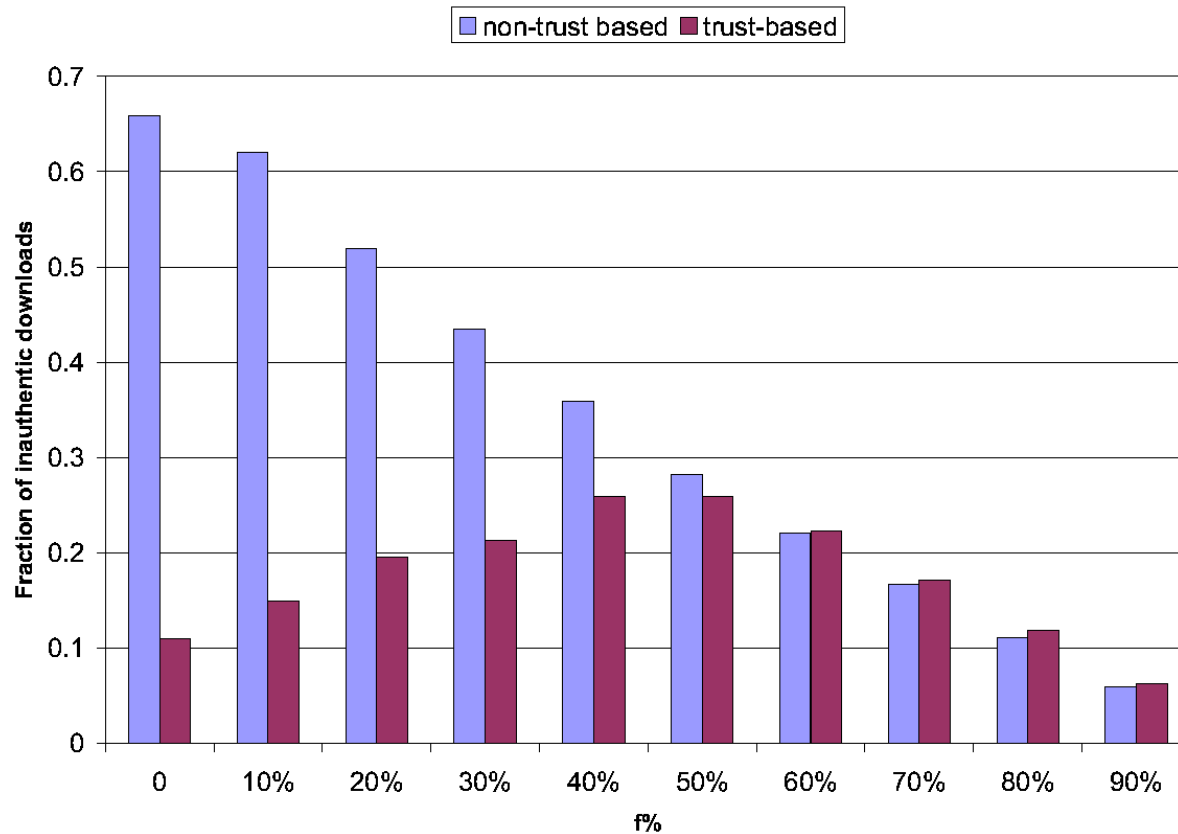
Malicious Individuals




Malicious Collective



Camouflaged Collective



Malicious Spies



Conclusion

- Eigentrust
 - Dramatically reduces number of inauthentic files on the network.
 - Robust to malicious peers.
 - Low overhead.
- Paper available at <http://www.stanford.edu/~sdkamvar/research.html>